

# Jersey Number Recognition

**Github repo:** <https://github.com/1morphed1/jersey-number-recognition>

**Best Model drive link:**

[https://drive.google.com/drive/folders/1sxm-wz\\_52PH0S751joqehYI4h2UXFpu0?usp=sharing](https://drive.google.com/drive/folders/1sxm-wz_52PH0S751joqehYI4h2UXFpu0?usp=sharing)

**Technical Report link:**

[https://github.com/1morphed1/jersey-number-recognition/blob/main/TECHNICAL\\_REPORT.md](https://github.com/1morphed1/jersey-number-recognition/blob/main/TECHNICAL_REPORT.md)


## Usage of AI:

For initial planning and validation I used claudie pro, deepseek, qwen, kimi, glm 4.6 and cursor agent, Then for finalizing the plan and designing the initial system claudie pro, deepseek and cursor was used. Cursor along with its mcp servers(context7, sequential thinking, exa search) was used all the time for checking on the correctness of my code. Cursor's memory bank was used to always have the full context of the project, and I also setup cursorrules so that the code format remains consistent overall. For debugging I used cursor agent, claudie pro and deepseek.


The cursor mcp servers integration helped with my knowledge gap and implementing boilerplate codes. Each script was run through over and over in claudie pro, deepseek and cursor , in order to maintain quality. Finally glm 4.6 used via CLine extension was used for file indexing , so that all the AI's which do not have knowledge bank like cursor can have full context to work with each time i open them. Also Claude Desktop along with mcp tools was used for final verification.The mcp tools were integrated from the site: <https://smithery.ai/>



The mermaid diagrams(flowcharts) were made with the help of file indexing capabilities of glm 4.6 agent mode, cursor agent mode and deepseek. Then visualized via the site:  
<https://mermaid.live/edit>



## Model Architecture:




## Tensorboard Logs:













## Process Flow

- Phase 1: Configuration and Setup
- Phase 2: Data Loading and Processing
- Phase 3: Model Architecture Creation
- Phase 4: Training Process
- Phase 5: Evaluation and Testing
- Phase 6: Model Analysis and Reporting









Now codes for individuals files will be shared below:

#### dataset.py:

```

# dataset.py

"""
Custom Dataset Loader for Jersey Number Recognition
Handles hierarchical directory structure:
Root/{Class}/{Jersey_Seq}/{Subdir}.jpg
"""

import os
import glob
import random
import torch
from torch.utils.data import Dataset
from PIL import Image, ImageEnhance
import numpy as np
from config import Config

```

```

class JerseySequenceDataset(Dataset):
    def __init__(self, root_dir, transform=None, seq_length=5,
                 allowed_classes=None, mode='train',
                 sampling_strategy='uniform'):
        """
        Args:
            root_dir (str): Root directory of dataset
            transform: torchvision transforms to apply
            seq_length (int): Number of frames to sample per sequence
            allowed_classes (list): List of class numbers to include (e.g.,
[4, 6, 46])
                                         If None, uses all available classes
            mode (str): 'train', 'val', or 'test' - for logging purposes
            sampling_strategy (str): 'uniform', 'random', or 'center'
        """
        self.root_dir = root_dir
        self.transform = transform
        self.seq_length = seq_length
        self.allowed_classes = allowed_classes
        self.mode = mode
        self.sampling_strategy = sampling_strategy
        self.samples = [] # List of (sequence_path, label_d1, label_d2,
class_name)

        self._crawl_dataset()

    def _parse_class_label(self, class_name):
        """
        Parse class folder name to digit labels.

        Args:
            class_name: Folder name
        Returns:
            (label_d1, label_d2): Tuple of digit labels
                - label_d1: Tens digit (0-9) or 10 for empty
                - label_d2: Units digit (0-9)
        """
        try:
            class_num = int(class_name)

```

```

        if class_num < 0 or class_num >= 100:
            raise ValueError(f"Class number out of range [0, 99]:
{class_num}")

        if class_num < 10:
            # Single digit: 0-9
            label_d1 = 10 # Empty tens place
            label_d2 = class_num
        else:
            # Two digits: 10-99
            label_d1 = class_num // 10
            label_d2 = class_num % 10

        return label_d1, label_d2

    except ValueError as e:
        print(f" Warning: Invalid class folder name '{class_name}':
{e}")
        return None, None

def _crawl_dataset(self):
    """
    Crawls the directory structure:
    Root/{Class}/{Jersey_Seq_ID}/{Label_Subdir}/.jpg
    """
    if not os.path.exists(self.root_dir):
        raise ValueError(f"Dataset path does not exist:
{self.root_dir}")

    # Get all class folders
    class_folders = [d for d in os.listdir(self.root_dir)
                    if os.path.isdir(os.path.join(self.root_dir, d))]

    # Filter by allowed classes if specified
    if self.allowed_classes is not None:
        class_folders = [c for c in class_folders
                        if c.isdigit() and int(c) in
self.allowed_classes]

```

```

    class_folders = sorted(class_folders, key=lambda x: int(x) if
x.isdigit() else 999)

    total_sequences = 0
    class_distribution = {}

    for class_name in class_folders:
        class_path = os.path.join(self.root_dir, class_name)

        # Parse labels using improved function
        label_d1, label_d2 = self._parse_class_label(class_name)

        # Skip invalid class names
        if label_d1 is None or label_d2 is None:
            continue

        # Iterate over Jersey/Sequence folders
        seq_folders = [f for f in os.listdir(class_path)
                      if os.path.isdir(os.path.join(class_path, f))]

        class_seq_count = 0

        for seq_id in seq_folders:
            seq_path = os.path.join(class_path, seq_id)

            # Iterate over internal label subdirectories
            # These are sub-sequences
            sub_dirs = [s for s in os.listdir(seq_path)
                        if os.path.isdir(os.path.join(seq_path, s))]

            for sub in sub_dirs:
                final_path = os.path.join(seq_path, sub)

                # Check if it contains images
                images = glob.glob(os.path.join(final_path, ".jpg")) +
\

                                glob.glob(os.path.join(final_path, ".png"))

                if len(images) >= 1: # At least 1 frame

```

```

                self.samples.append((final_path, label_d1,
label_d2, class_name))
                class_seq_count += 1

        if class_seq_count > 0:
            class_distribution[class_name] = class_seq_count
            total_sequences += class_seq_count

    # Print dataset statistics
    print(f"\n{'='*60}")
    print(f"📊 {self.mode.upper()} Dataset Loaded")
    print(f"{'='*60}")
    print(f"Total sequences: {total_sequences}")
    print(f"Classes found: {len(class_distribution)}")
    print("\nClass distribution:")
    for cls, count in sorted(class_distribution.items(), key=lambda x:
int(x[0])):
        print(f"  Class {cls:>2s}: {count:>4d} sequences")
    print(f"{'='*60}\n")

def _sample_frames(self, image_files):
    """
    Sample frames using different strategies, prioritizing anchor
    frames.

    Args:
        image_files: List of image file paths

    Returns:
        indices: List of frame indices to use
    """
    num_frames = len(image_files)
    indices = []

    # Find anchor frames
    anchor_indices = [i for i, f in enumerate(image_files) if '_anchor'
in f]

    # Always include at least one anchor frame if available
    if anchor_indices:

```

```

        # Use the first anchor as frame 0
        indices.append(anchor_indices[0])
        # Remove it from candidate pool to avoid duplication
        remaining_indices = [i for i in range(num_frames) if i not in
anchor_indices]
    else:
        # No anchors - sample from all frames
        remaining_indices = list(range(num_frames))

    # Sample remaining frames based on strategy
    remaining_needed = self.seq_length - len(indices)

    if remaining_needed > 0:
        if self.sampling_strategy == 'uniform':
            if len(remaining_indices) >= remaining_needed:
                step = len(remaining_indices) / remaining_needed
                sampled = [int(i * step) for i in
range(remaining_needed)]
                indices.extend([remaining_indices[i] for i in sampled])
            else:
                indices.extend(remaining_indices)
                while len(indices) < self.seq_length:
                    indices.append(remaining_indices[-1])

        elif self.sampling_strategy == 'random':
            if len(remaining_indices) >= remaining_needed:
                sampled = np.random.choice(remaining_indices,
remaining_needed, replace=False)
                indices.extend(sampled.tolist())
            else:
                indices.extend(remaining_indices)
                while len(indices) < self.seq_length:
                    indices.append(np.random.choice(remaining_indices))

        elif self.sampling_strategy == 'center':
            center = len(remaining_indices) // 2
            half_seq = remaining_needed // 2
            start = max(0, center - half_seq)
            end = min(len(remaining_indices), start + remaining_needed)
            sampled = remaining_indices[start:end]

```

```

        indices.extend(sampled)
        while len(indices) < self.seq_length:
            indices.append(sampled[-1])

    else:
        raise ValueError(f"Unknown sampling strategy:
{self.sampling_strategy}")

    # Ensure we have exactly seq_length frames
    indices = indices[:self.seq_length]
    while len(indices) < self.seq_length:
        indices.append(indices[-1])

    return indices

def __len__(self):
    return len(self.samples)

def _enhance_image(self, img):
    """
    Enhance image quality: sharpen and optionally upscale.

    Args:
        img: PIL.Image

    Returns:
        Enhanced PIL.Image
    """
    # Option 1: Sharpen using PIL (lightweight, no extra deps)
    enhancer = ImageEnhance.Sharpness(img)
    img = enhancer.enhance(2.0)  # Increase sharpness by 2x

    # Optional: Upscale if image is very small (e.g., < 64px)
    min_size = 64
    if img.width < min_size or img.height < min_size:
        # Use LANCZOS for best quality upsampling
        scale_factor = max(min_size / img.width, min_size / img.height)
        new_size = (int(img.width * scale_factor), int(img.height *
scale_factor))
        img = img.resize(new_size, Image.LANCZOS)

```

```
    return img

def __getitem__(self, idx):
    seq_path, d1, d2, class_name = self.samples[idx]

    # Get all image files
    image_files = sorted(
        glob.glob(os.path.join(seq_path, ".jpg")) +
        glob.glob(os.path.join(seq_path, ".png")))
    )

    # Sample frames using chosen strategy
    indices = self._sample_frames(image_files)

    # Load and Transform Images
    frames = []
    for i in indices:
        img_path = image_files[i]
        try:
            img = Image.open(img_path).convert('RGB')

            # ENHANCE IMAGE BEFORE TRANSFORM
            img = self._enhance_image(img)

            if self.transform:
                img = self.transform(img)
            else:
                # Always resize + normalize for consistency
                from torchvision import transforms
                default_transform = transforms.Compose([
                    transforms.Resize((Config.IMG_SIZE,
Config.IMG_SIZE)),
                    transforms.ToTensor(),
                    transforms.Normalize(mean=[0.485, 0.456, 0.406],
std=[0.229, 0.224, 0.225])
                ])
            img = default_transform(img)
            frames.append(img)
        except Exception as e:
```

```

        print(f" Warning: Failed to load {img_path}: {e}")
        if self.transform:
            black_frame = torch.zeros(3, Config.IMG_SIZE,
Config.IMG_SIZE)
        else:
            black_frame = torch.zeros(3, 224, 224)
        frames.append(black_frame)

    # Stack into tensor: [Seq_Len, Channels, Height, Width]
    seq_tensor = torch.stack(frames)

    return seq_tensor, torch.tensor(d1, dtype=torch.long),
torch.tensor(d2, dtype=torch.long)

    def get_class_weights(self, method='balanced',
suppress_empty_weight=True):
        """
        Calculate class weights for handling imbalanced data.

    Args:
        method: 'balanced' or 'effective_number'
        suppress_empty_weight: If True, heavily suppress the weight for
D1=10 (empty)

    Returns:
        (weights_d1, weights_d2): Tuple of weight tensors
        """
        from collections import Counter

        d1_labels = [s[1] for s in self.samples]
        d2_labels = [s[2] for s in self.samples]

        d1_counts = Counter(d1_labels)
        d2_counts = Counter(d2_labels)

        print("\n Label Distribution:")
        print(f"D1 counts: {dict(sorted(d1_counts.items()))}")
        print(f"D2 counts: {dict(sorted(d2_counts.items()))}")

    # Initialize weights

```

```

d1_weights = torch.ones(Config.NUM_DIGIT_CLASSES)
d2_weights = torch.ones(Config.NUM_DIGIT_CLASSES)

total_samples = len(self.samples)

if method == 'balanced':
    # Inverse frequency weighting
    for digit in range(Config.NUM_DIGIT_CLASSES):
        if digit in d1_counts and d1_counts[digit] > 0:
            d1_weights[digit] = total_samples / (d1_counts[digit])
Config.NUM_DIGIT_CLASSES)
        else:
            # Missing class: give very low weight
            d1_weights[digit] = 0.001

    for digit in range(Config.NUM_DIGIT_CLASSES):
        if digit in d2_counts and d2_counts[digit] > 0:
            d2_weights[digit] = total_samples / (d2_counts[digit])
Config.NUM_DIGIT_CLASSES)
        else:
            d2_weights[digit] = 0.001

elif method == 'effective_number':
    # Effective number of samples
    beta = 0.9999

    for digit in range(Config.NUM_DIGIT_CLASSES):
        if digit in d1_counts and d1_counts[digit] > 0:
            effective_num = (1.0 - beta * d1_counts[digit]) / (1.0 -
beta)
            d1_weights[digit] = 1.0 / effective_num
        else:
            d1_weights[digit] = 0.001

    for digit in range(Config.NUM_DIGIT_CLASSES):
        if digit in d2_counts and d2_counts[digit] > 0:
            effective_num = (1.0 - beta * d2_counts[digit]) / (1.0 -
beta)
            d2_weights[digit] = 1.0 / effective_num
        else:

```

```

d2_weights[digit] = 0.001

# Manually suppress the "empty" class (10) for D1
if suppress_empty_weight and 10 in d1_counts:
    print(f"\n🔧 Suppressing D1 'empty' class weight by 100x")
    d1_weights[10] = d1_weights[10] * 0.01

else:
    raise ValueError(f"Unknown class weight method: {method}")

d1_weights = torch.sqrt(d1_weights)
d2_weights = torch.sqrt(d2_weights)

# Print weight distribution with better formatting
print(f"\n Class Weights (method={method}):")
print("\nD1 weights (Tens digit):")
for i in range(Config.NUM_DIGIT_CLASSES):
    label = "Empty" if i == 10 else str(i)
    count = d1_counts.get(i, 0)
    print(f" {label:>5s} (count: {count:>6d}): weight = "
{d1_weights[i]:.6f})

print("\nD2 weights (Units digit):")
for i in range(Config.NUM_DIGIT_CLASSES):
    label = "Empty" if i == 10 else str(i)
    count = d2_counts.get(i, 0)
    if count > 0:
        print(f" {label:>5s} (count: {count:>6d}): weight = "
{d2_weights[i]:.6f})

return d1_weights, d2_weights

def set_transform(self, transform):
    """Change the transform"""
    self.transform = transform


# Test the dataset
if __name__ == "__main__":
    from torchvision import transforms

```

```

from config import Config, validate_config

validate_config()

# Define transforms
transform = transforms.Compose([
    transforms.Resize((Config.IMG_SIZE, Config.IMG_SIZE)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225])
])

# Create dataset
dataset = JerseySequenceDataset(
    root_dir=Config.DATA_ROOT,
    transform=transform,
    seq_length=Config.SEQ_LENGTH,
    allowed_classes=Config.TRAIN_CLASSES,
    mode='train',
    sampling_strategy='uniform'
)

print("\n Dataset created successfully!")
print(f"Total samples: {len(dataset)}")

# Test loading one sample
if len(dataset) > 0:
    seq, d1, d2 = dataset[0]
    print(f"\nSample shape: {seq.shape}")
    print(f"Label D1 (tens): {d1.item()}")
    print(f"Label D2 (units): {d2.item()}")

    # Test class weights
    w1, w2 = dataset.get_class_weights(method='balanced')
    print("\n Class weights calculated!")

```

This is a custom PyTorch Dataset class for loading jersey number recognition data from video sequences.

## 1. WHAT THE DATASET DOES:

Loads sequences of images showing jersey numbers (like sports jerseys with numbers 0-99)

Each sequence shows the same jersey number from multiple frames

The goal is to train a model to recognize two-digit numbers

## 2. DIRECTORY STRUCTURE IT EXPECTS:

Root/Class/Sequence\_ID/Subfolder/.jpg

Example:

Data/4/Seq\_001/labeled\_001/frame1.jpg

Data/4/Seq\_001/labeled\_001/frame2.jpg

...

Data/46/Seq\_003/labeled\_002/frame1.jpg

## 3. KEY FEATURES OF THE DATASET CLASS:

Label Parsing:

Jersey numbers are parsed into two digits

Single-digit numbers (0-9): first digit is "empty" (label 10), second digit is the number

Two-digit numbers (10-99): first digit is tens place, second digit is units place

Example: "4" becomes (10, 4), "46" becomes (4, 6)

Frame Sampling Strategies:

'uniform': evenly spaced frames from the sequence

'random': random frames

'center': frames from middle of sequence

Prioritizes "anchor" frames (frames with '\_anchor' in filename) as first frame

Image Enhancement:

Sharpens images to improve clarity

Optionally upscales small images (<64px) for better quality

Applies standard transforms (resize, normalize)

Class Weight Calculation:

Calculates weights for imbalanced datasets

Two methods: 'balanced' (inverse frequency) and 'effective\_number'

Can suppress weight for "empty" tens digit to avoid bias

4. WORKFLOW:

Initialization (\_crawl\_dataset):

Scans directory structure

Counts sequences per class

Filters by allowed classes if specified

Prints statistics about dataset

Getting an item (getitem):

Takes a sequence path

Samples frames using chosen strategy

Loads and enhances each image

Applies transforms

Returns: tensor of shape [Seq\_Length, Channels, Height, Width], first digit label, second digit label

5. OUTPUT FORMAT:

For each sample returns:

seq\_tensor: [5, 3, 224, 224] (if seq\_length=5, image size=224)

d1\_label: tensor with first digit (0-9 or 10 for empty)

d2\_label: tensor with second digit (0-9)

## 6. USE CASE:

This dataset is designed for training models that:

Take multiple frames of a jersey number as input

Predict two separate digits (tens and units place)

Handle single-digit numbers (like "7") and two-digit numbers (like "23")

Work with potentially imbalanced class distributions

## 7. TESTING:

The bottom section tests the dataset by:

Creating a transform (resize, normalize)

Instantiating the dataset

Loading one sample

Calculating class weights

The main strength is handling sequential data with hierarchical directory structures and providing multiple sampling strategies for frames.

### config.py:

```
# config.py

"""
Configuration file for Jersey Number Recognition
Optimized for your actual dataset
"""

import os
import warnings
from dotenv import load_dotenv

load_dotenv() # Load environment variables from .env file if present

class Config:
    # =====
```

```
# DATA CONFIGURATION
# =====
# Use environment variable
DATA_ROOT = os.environ.get('DATA_ROOT')

# =====
# TRAINING/TEST SPLIT
# =====

TRAIN_CLASSES = [4, 6, 8, 9, 49, 66, 89]
TEST_CLASSES = [48, 64, 88]

# Validation split (from training data)
VAL_SPLIT = 0.2

# =====
# MODEL CONFIGURATION
# =====

NUM_DIGIT_CLASSES = 11 # 0-9 + empty = 11
HIDDEN_DIM = 256
IMG_SIZE = 64
SEQ_LENGTH = 5

# Model architecture options
BACKBONE = 'resnet18'
USE_SPATIAL_ATTENTION = True
USE_TEMPORAL_ATTENTION = True
BIDIRECTIONAL_LSTM = True

# =====
# TRAINING CONFIGURATION
# =====

BATCH_SIZE = 16
NUM_EPOCHS = 30
LEARNING_RATE = 0.001
WEIGHT_DECAY = 1e-4
DROPOUT = 0.3

# Gradient clipping
GRADIENT_CLIP_VALUE = 1.0
```

```
# Learning rate scheduler
USE_SCHEDULER = True
SCHEDULER_STEP_SIZE = 10
SCHEDULER_GAMMA = 0.5

# Early stopping
EARLY_STOPPING_PATIENCE = 10
EARLY_STOPPING_MIN_DELTA = 0.001

# Mixed precision training
USE_MIXED_PRECISION = True

# Class weight method
CLASS_WEIGHT_METHOD = 'effective_number' # 'balanced' or
'effective_number'

# =====
# DATA AUGMENTATION
# =====
# Safe augmentations that don't change digit identity
AUG_ROTATION_DEGREES = 5
AUG_TRANSLATE = (0.05, 0.05)
AUG_SCALE = (0.95, 1.05)
AUG_BRIGHTNESS = 0.2
AUG_CONTRAST = 0.2

# Temporal sampling strategy
TEMPORAL_SAMPLING = 'uniform' # 'uniform', 'random', 'center'

# =====
# SYSTEM CONFIGURATION
# =====
NUM_WORKERS = 4
DEVICE = 'cuda'
SEED = 42
PIN_MEMORY = True

# Confidence thresholds for zero-shot jersey number reconstruction
```

```

D1_CONF_THRESHOLD = 0.10    # Tens digit confidence threshold
D2_CONF_THRESHOLD = 0.5     # Units digit confidence threshold

# =====
# OUTPUT PATHS
# =====
CHECKPOINT_DIR = 'checkpoints'
RESULTS_DIR = 'results'
LOGS_DIR = 'logs'

@staticmethod
def create_dirs():
    """Create all necessary output directories"""
    os.makedirs(Config.CHECKPOINT_DIR, exist_ok=True)
    os.makedirs(Config.RESULTS_DIR, exist_ok=True)
    os.makedirs(Config.LOGS_DIR, exist_ok=True)
    os.makedirs(os.path.join(Config.RESULTS_DIR, 'confusion_matrices'),
exist_ok=True)
    os.makedirs(os.path.join(Config.RESULTS_DIR, 'predictions'),
exist_ok=True)
    os.makedirs(os.path.join(Config.RESULTS_DIR, 'visualizations'),
exist_ok=True)

@staticmethod
def validate():
    """Validate configuration parameters"""
    errors = []
    warnings_list = []

    # Check data root exists
    if not os.path.exists(Config.DATA_ROOT):
        errors.append(f"Data root does not exist: {Config.DATA_ROOT}")

    # Check train/test split
    if Config.TRAIN_CLASSES is not None and Config.TEST_CLASSES is not
None:
        overlap = set(Config.TRAIN_CLASSES) & set(Config.TEST_CLASSES)
        if overlap:
            errors.append(f"Train and test classes overlap: {overlap}")

```

```
# Validate hyperparameters
if Config.BATCH_SIZE < 1:
    errors.append("Batch size must be >= 1")

if not (32 <= Config.IMG_SIZE <= 512):
    warnings_list.append(f"IMG_SIZE={Config.IMG_SIZE} is unusual
(recommended: 128-256)")

if Config.SEQ_LENGTH < 1:
    errors.append("SEQ_LENGTH must be >= 1")

if not (0.0 < Config.VAL_SPLIT < 1.0):
    errors.append("VAL_SPLIT must be between 0 and 1")

if not (0.0 <= Config.DROPOUT < 1.0):
    errors.append("DROPOUT must be between 0 and 1")

# Check GPU availability
if Config.DEVICE == 'cuda':
    try:
        import torch
        if not torch.cuda.is_available():
            warnings_list.append("CUDA requested but not available,
will fall back to CPU")
            Config.DEVICE = 'cpu'
    except ImportError:
        errors.append("PyTorch not installed")

# Print errors
if errors:
    print("\n CONFIGURATION ERRORS:")
    for error in errors:
        print(f"  • {error}")
    raise ValueError("Configuration validation failed")

# Print warnings
if warnings_list:
    print("\n CONFIGURATION WARNINGS:")
    for warning in warnings_list:
        print(f"  • {warning}")
```

```
    return True

def validate_config():
    """Check if configuration is valid and analyze digit coverage"""
    print("\n" + "="*70)
    print("VALIDATING CONFIGURATION")
    print("="*70)

    # Validate parameters
    Config.validate()

    # Validate digit coverage
    print("\n🔍 Analyzing Digit Coverage:")

    if Config.TRAIN_CLASSES is None:
        print(" No specific training classes defined")
        return

    train_tens = set()
    train_units = set()

    for cls in Config.TRAIN_CLASSES:
        if cls < 10:
            train_units.add(cls)
        else:
            train_tens.add(cls // 10)
            train_units.add(cls % 10)

    print(f" Training tens digits: {sorted(train_tens)} + empty")
    print(f" Training units digits: {sorted(train_units)}")

    # Check generalization to test classes
    if Config.TEST_CLASSES:
        print(f"\n Checking generalization to test classes:")
        can_generalize_all = True

        for test_cls in Config.TEST_CLASSES:
            if test_cls < 10:
```

```

        if test_cls in train_units:
            print(f" Class {test_cls}: covered (single digit)")
        else:
            print(f" Class {test_cls}: NOT covered")
            can_generalize_all = False
    else:
        tens = test_cls // 10
        units = test_cls % 10

        tens_covered = tens in train_tens
        units_covered = units in train_units

        if tens_covered and units_covered:
            print(f" Class {test_cls} = {tens}{units}: covered")
        else:
            missing = []
            if not tens_covered:
                missing.append(f"tens={tens}")
            if not units_covered:
                missing.append(f"units={units}")
            print(f" Class {test_cls}: NOT covered ({',
'.join(missing)})")
            can_generalize_all = False

    if can_generalize_all:
        print("\n Model should generalize to all test classes!")
    else:
        print("\n WARNING: Model may struggle with some test classes")

print("\n" + "="*70)
print("CONFIGURATION SUMMARY")
print("=="*70)
print(f" Data root: {Config.DATA_ROOT}")
print(f" Train classes: {Config.TRAIN_CLASSES}")
print(f" Test classes: {Config.TEST_CLASSES}")
print(f" Validation split: {Config.VAL_SPLIT100:.0f}%")
print(f" Backbone: {Config.BACKBONE}")
print(f" Image size: {Config.IMG_SIZE}x{Config.IMG_SIZE}")
print(f" Sequence length: {ConfigSEQ_LENGTH} frames")
print(f" Batch size: {Config.BATCH_SIZE}")

```

```

print(f" Epochs: {Config.NUM_EPOCHS}")
print(f" Learning rate: {Config.LEARNING_RATE}")
print(f" Checkpoints: {Config.CHECKPOINT_DIR}/")
print(f" Results: {Config.RESULTS_DIR}/")
print("=="70 + "\n")

if __name__ == "__main__":
    validate_config()

```

This is a configuration file that sets all the parameters for training a jersey number recognition model.

#### WHAT THIS FILE DOES:

##### 1. Loads environment variables

Reads a .env file if it exists

Gets the DATA\_ROOT path from environment variable (so you don't hardcode paths)

##### 2. Defines configuration parameters in categories:

#### DATA CONFIGURATION:

DATA\_ROOT: Where your dataset is located

TRAIN\_CLASSES: Which jersey numbers to use for training (7 numbers: 4, 6, 8, 9, 49, 66, 89)

TEST\_CLASSES: Which jersey numbers to use for testing (3 numbers: 48, 64, 88)

VAL\_SPLIT: 20% of training data used for validation

#### MODEL CONFIGURATION:

NUM\_DIGIT\_CLASSES: 11 (digits 0-9 + "empty" for no tens digit)

IMG\_SIZE: 64x64 pixels (small images)

SEQ\_LENGTH: 5 frames per sequence

**BACKBONE:** Uses ResNet18 as the base model

Uses both spatial and temporal attention

Uses bidirectional LSTM

#### TRAINING CONFIGURATION:

BATCH\_SIZE: 16 images per batch

NUM\_EPOCHS: 30 training cycles

LEARNING\_RATE: 0.001

Uses mixed precision training (faster, less memory)

Uses class weights to handle imbalanced data

#### DATA AUGMENTATION:

Small rotations (5 degrees)

Small translations

Brightness/contrast adjustments

All safe augmentations that don't change the digit identity

#### SYSTEM CONFIGURATION:

NUM\_WORKERS: 4 parallel data loaders

Uses GPU (cuda) if available

Fixed random seed (42) for reproducibility

#### OUTPUT PATHS:

Creates directories for checkpoints, results, logs

#### 3. Key features of this config:

##### Class Coverage Analysis:

The validate\_config function checks if the training data covers all digits needed for test numbers:

Training has digits: 4, 6, 8, 9, 4, 6, 6, 8, 9 (from 49, 66, 89)

Testing has digits: 4, 8, 6, 4, 8, 8 (from 48, 64, 88)

Since test digits (4, 6, 8) are all in training, the model should generalize

Example:

Test number 48:

Tens digit 4: Seen in training (from 49)

Units digit 8: Seen in training (from 8 and 89)

Both digits are covered, so model can predict 48

4. Configuration Validation:

The validate() method checks for:

Data path exists

No overlap between train and test classes

Hyperparameters are reasonable

GPU availability

5. Directory Creation:

The create\_dirs() method makes sure all necessary folders exist for saving:

Model checkpoints

Results

Logs

Confusion matrices

Visualizations

6. Main Idea:

This config separates train/test by jersey numbers:

Train on some numbers (4, 6, 8, 9, 49, 66, 89)

Test on different numbers (48, 64, 88)

But all individual digits (4, 6, 8, 9) appear in both sets

This tests if the model learns to recognize individual digits, not just memorizing specific jersey numbers

The validation function shows if this approach will work by checking digit coverage between train and test sets.

### [model.py](#)

```
# model.py

"""
Enhanced Two-Head Temporal Jersey Recognition Model (CRNN)
Architecture: ResNet18/34 (CNN) + Bidirectional LSTM + Temporal Attention
+ Dual Classification Heads

Key Features:
- Temporal attention mechanism (no information bottleneck)
- Bidirectional LSTM for better context
- Spatial attention for focusing on jersey region
- Flexible backbone selection
- Better classification heads with batch normalization
- Support for variable-length sequences with padding masks
- Gradient clipping support
"""

import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import models
from typing import Tuple, Optional

class SpatialAttention(nn.Module):
```

```

"""
Spatial attention module to focus on jersey number regions
Generates attention weights across spatial dimensions
"""

def __init__(self, in_channels):
    super(SpatialAttention, self).__init__()
    self.conv = nn.Sequential(
        nn.Conv2d(in_channels, in_channels // 8, kernel_size=1),
        nn.BatchNorm2d(in_channels // 8),
        nn.ReLU(inplace=True),
        nn.Conv2d(in_channels // 8, 1, kernel_size=1),
        nn.Sigmoid()
    )

    def forward(self, x):
        """
        Args:
            x: [Batch, Channels, H, W]
        Returns:
            Attention-weighted features
        """
        attention = self.conv(x)  # [Batch, 1, H, W]
        return x * attention


class TemporalAttention(nn.Module):
    """
    Temporal attention mechanism to weight different frames
    Learns which frames are most informative for digit recognition
    """

    def __init__(self, hidden_dim):
        super(TemporalAttention, self).__init__()
        self.attention = nn.Sequential(
            nn.Linear(hidden_dim, hidden_dim // 2),
            nn.Tanh(),
            nn.Linear(hidden_dim // 2, 1)
        )

    def forward(self, lstm_out, mask=None):
        """
"""

```

```

Args:
    lstm_out: [Batch, Seq, Hidden]
    mask: [Batch, Seq] - Optional padding mask (1 for valid, 0 for
padding)

Returns:
    context: [Batch, Hidden] - Weighted sum of temporal features
    weights: [Batch, Seq] - Attention weights
"""

# Compute attention scores
scores = self.attention(lstm_out) # [Batch, Seq, 1]
scores = scores.squeeze(-1) # [Batch, Seq]

# Apply mask if provided
if mask is not None:
    scores = scores.masked_fill(mask == 0, -1e9)

# Softmax to get attention weights
weights = F.softmax(scores, dim=1) # [Batch, Seq]

# Weighted sum
context = torch.bmm(weights.unsqueeze(1), lstm_out) # [Batch, 1,
Hidden]
context = context.squeeze(1) # [Batch, Hidden]

return context, weights


class EnhancedDigitHead(nn.Module):
"""

Enhanced classification head with batch normalization and residual
connection
"""

def __init__(self, input_dim, num_classes, dropout=0.3):
    super(EnhancedDigitHead, self).__init__()

    hidden_dim = input_dim // 2

    self.fc1 = nn.Linear(input_dim, hidden_dim)
    self.bn1 = nn.BatchNorm1d(hidden_dim)
    self.dropout1 = nn.Dropout(dropout)

```

```

    self.fc2 = nn.Linear(hidden_dim, hidden_dim)
    self.bn2 = nn.BatchNorm1d(hidden_dim)
    self.dropout2 = nn.Dropout(dropout)

    self.fc_out = nn.Linear(hidden_dim, num_classes)

    # Residual connection projection if needed
    self.shortcut = nn.Linear(input_dim, hidden_dim) if input_dim != hidden_dim else None

def forward(self, x):
    """
    Args:
        x: [Batch, input_dim]
    Returns:
        logits: [Batch, num_classes]
    """
    identity = x

    # First block
    out = self.fc1(x)
    out = self.bn1(out)
    out = F.relu(out)
    out = self.dropout1(out)

    # Second block with residual
    if self.shortcut is not None:
        identity = self.shortcut(identity)

        out = self.fc2(out)
        out = self.bn2(out)
        out = out + identity # Residual connection
        out = F.relu(out)
        out = self.dropout2(out)

    # Output
    logits = self.fc_out(out)

    return logits

```

```

class JerseyTemporalNet(nn.Module):
    """
    Enhanced Two-Head Temporal Jersey Recognition Model
    """

    def __init__(
        self,
        num_classes=11,
        hidden_dim=256,
        dropout=0.3,
        backbone='resnet18',
        use_spatial_attention=True,
        use_temporal_attention=True,
        bidirectional=True
    ):
        """
        Args:
            num_classes (int): Number of digit classes (0-9 + empty = 11)
            hidden_dim (int): LSTM hidden dimension
            dropout (float): Dropout probability
            backbone (str): CNN backbone ('resnet18', 'resnet34',
            'resnet50')
            use_spatial_attention (bool): Whether to use spatial attention
            use_temporal_attention (bool): Whether to use temporal
            attention
            bidirectional (bool): Whether to use bidirectional LSTM
        """
        super(JerseyTemporalNet, self).__init__()

        self.use_spatial_attention = use_spatial_attention
        self.use_temporal_attention = use_temporal_attention
        self.bidirectional = bidirectional

        # =====
        # A. SPATIAL FEATURE EXTRACTOR (CNN BACKBONE)
        # =====
        self.backbone_name = backbone
        self.backbone, self.feature_dim = self._build_backbone(backbone)

```

```

# Spatial attention module
if use_spatial_attention:
    # Get the number of channels before global pooling
    # For ResNet: layer4 outputs different channels per variant
    if 'resnet18' in backbone or 'resnet34' in backbone:
        spatial_channels = 512
    else:
        spatial_channels = 2048
    self.spatial_attention = SpatialAttention(spatial_channels)

# =====
# B. TEMPORAL AGGREGATOR (LSTM)
# =====

self.lstm = nn.LSTM(
    input_size=self.feature_dim,
    hidden_size=hidden_dim,
    num_layers=2,
    batch_first=True,
    dropout=dropout if dropout > 0 else 0,
    bidirectional=bidirectional
)

# Effective hidden dimension after LSTM
lstm_output_dim = hidden_dim * 2 if bidirectional else hidden_dim

# Temporal attention module
if use_temporal_attention:
    self.temporal_attention = TemporalAttention(lstm_output_dim)
    self.aggregation_method = 'attention'
else:
    self.aggregation_method = 'mean' # Fallback to mean pooling

# =====
# C. CLASSIFICATION HEADS
# =====

self.dropout = nn.Dropout(dropout)

# Head 1: Tens Place Digit (0-9, or 10 for "empty")
self.head_digit1 = EnhancedDigitHead(lstm_output_dim, num_classes,
dropout)

```

```

# Head 2: Units Place Digit (0-9)
self.head_digit2 = EnhancedDigitHead(lstm_output_dim, num_classes,
dropout)

# Initialize weights
self._initialize_weights()

def _freeze_backbone_layers(self, backbone, freeze_percentage=0.6):
    """
    Freeze backbone layers intelligently.

    Args:
        backbone: The CNN backbone module
        freeze_percentage: Percentage of layers to freeze (0.0 to 1.0)

    Returns:
        Number of frozen parameters, total parameters
    """
    # Get all parameters with their names
    params = list(backbone.named_parameters())
    total_layers = len(params)
    freeze_until_idx = int(total_layers * freeze_percentage)

    frozen_params = 0
    total_params = 0

    print(f"\n Freezing {freeze_percentage*100:.0f}% of backbone layers\n({freeze_until_idx}/{total_layers})")

    for idx, (name, param) in enumerate(params):
        total_params += param.numel()

        if idx < freeze_until_idx:
            param.requires_grad = False
            frozen_params += param.numel()
            if idx < 3:
                print(f" Frozen: {name} ({param.numel():,} params)")
        else:
            param.requires_grad = True

```

```

        if idx == freeze_until_idx: # Print first unfrozen
            print(f" ... ({freeze_until_idx - 3} more frozen
layers)")

            print(f" Trainable: {name} ({param.numel():,} params)")

        elif idx < freeze_until_idx + 2:
            print(f" Trainable: {name} ({param.numel():,} params)")

    print(f"\n  Total: {frozen_params:,} / {total_params:,} frozen
({frozen_params/total_params*100:.1f}%)")

    return frozen_params, total_params

def _build_backbone(self, backbone_name):
    """Build and configure the CNN backbone"""
    if backbone_name == 'resnet18':
        resnet = models.resnet18(pretrained=True)
    elif backbone_name == 'resnet34':
        resnet = models.resnet34(pretrained=True)
    elif backbone_name == 'resnet50':
        resnet = models.resnet50(pretrained=True)
    else:
        raise ValueError(f"Unsupported backbone: {backbone_name}")

    feature_dim = resnet.fc.in_features

    # Remove final FC and global pooling
    if self.use_spatial_attention:
        # Keep layer4 for spatial attention, remove avgpool and fc
        backbone = nn.Sequential(list(resnet.children())[:-2])
    else:
        # Remove only the fc layer
        backbone = nn.Sequential(list(resnet.children())[:-1])

    # Freeze layers intelligently
    self._freeze_backbone_layers(backbone, freeze_percentage=0.6)

    return backbone, feature_dim

def _initialize_weights(self):
    """Initialize weights for new layers"""

```

```

for m in self.modules():
    if isinstance(m, nn.Linear):
        nn.init.xavier_uniform_(m.weight)
        if m.bias is not None:
            nn.init.constant_(m.bias, 0)
    elif isinstance(m, nn.LSTM):
        for name, param in m.named_parameters():
            if 'weight_ih' in name:
                nn.init.xavier_uniform_(param.data)
            elif 'weight_hh' in name:
                nn.init.orthogonal_(param.data)
            elif 'bias' in name:
                nn.init.constant_(param.data, 0)

def unfreeze_backbone(self, percentage=0.3):
    """
    Gradually unfreeze more backbone layers
    """

    Args:
        percentage (float): Percentage of layers to unfreeze (0.0 to 1.0)
    """
    params = list(self.backbone.named_parameters())
    num_to_unfreeze = int(len(params) * percentage)

    print(f"\n Unfreezing {percentage*100:.0f}% of backbone layers\n({num_to_unfreeze} layers):")

    for idx, (name, param) in enumerate(params[-num_to_unfreeze:]):
        param.requires_grad = True
        if idx < 3: # Print first 3
            print(f" Unfrozen: {name}")

def forward(self, x, mask=None):
    """
    Forward pass
    """

    Args:
        x: Input tensor [Batch, Seq_Length, Channels, Height, Width]

```

```
    mask: Optional padding mask [Batch, Seq_Length] (1 for valid, 0
for padding)

    Returns:
        d1_logits: Logits for tens place [Batch, num_classes]
        d2_logits: Logits for units place [Batch, num_classes]
        attention_weights: Temporal attention weights (if enabled)
[Batch, Seq]

"""
batch_size, seq_length, c, h, w = x.size()

# =====
# Step 1: Extract spatial features from each frame
# =====
x = x.view(batch_size  seq_length, c, h, w)

# Pass through CNN backbone
features = self.backbone(x)

# Apply spatial attention if enabled
if self.use_spatial_attention and features.dim() == 4:
    features = self.spatial_attention(features)
    # Global average pooling
    features = F.adaptive_avg_pool2d(features, (1, 1))

# Flatten spatial dimensions
features = features.view(batch_size  seq_length, -1)

# Reshape back to sequence
features = features.view(batch_size, seq_length, -1)

# =====
# Step 2: Temporal modeling with LSTM
# =====
if mask is not None:
    # Pack padded sequence for efficiency
    lengths = mask.sum(dim=1).cpu()
    packed = nn.utils.rnn.pack_padded_sequence(
        features, lengths, batch_first=True, enforce_sorted=False
    )
```

```

        lstm_out, (h_n, c_n) = self.lstm(packed)
        lstm_out, _ = nn.utils.rnn.pad_packed_sequence(lstm_out,
batch_first=True)
    else:
        lstm_out, (h_n, c_n) = self.lstm(features)

        # =====
        # Step 3: Temporal aggregation
        # =====
        attention_weights = None

        if self.use_temporal_attention:
            # Use attention mechanism
            aggregated, attention_weights =
self.temporal_attention(lstm_out, mask)
        else:
            # Fallback to mean pooling
            if mask is not None:
                # Masked mean
                mask_expanded = mask.unsqueeze(-1).expand_as(lstm_out)
                sum_out = (lstm_out * mask_expanded).sum(dim=1)
                lengths = mask.sum(dim=1, keepdim=True).float()
                aggregated = sum_out / lengths
            else:
                aggregated = lstm_out.mean(dim=1)

        # Apply dropout
        aggregated = self.dropout(aggregated)

        # =====
        # Step 4: Digit classification
        # =====
        d1_logits = self.head_digit1(aggregated)
        d2_logits = self.head_digit2(aggregated)

        if attention_weights is not None:
            return d1_logits, d2_logits, attention_weights
        else:
            return d1_logits, d2_logits

```

```
def predict(self, x, mask=None, return_confidence=True):
    """
    Predict digit labels (for inference)

    Args:
        x: Input tensor [Batch, Seq_Length, Channels, Height, Width]
        mask: Optional padding mask [Batch, Seq_Length]
        return_confidence: Whether to return confidence scores

    Returns:
        d1_pred: Predicted tens digit [Batch]
        d2_pred: Predicted units digit [Batch]
        confidence: Average confidence scores [Batch] (if
return_confidence=True)
        attention_weights: Temporal attention weights [Batch, Seq] (if
enabled)
    """
    self.eval()
    with torch.no_grad():
        forward_out = self.forward(x, mask)

        if len(forward_out) == 3:
            d1_logits, d2_logits, attention_weights = forward_out
        else:
            d1_logits, d2_logits = forward_out
            attention_weights = None

        # Get probabilities
        d1_probs = F.softmax(d1_logits, dim=1)
        d2_probs = F.softmax(d2_logits, dim=1)

        # Get predictions
        d1_pred = torch.argmax(d1_probs, dim=1)
        d2_pred = torch.argmax(d2_probs, dim=1)

        if return_confidence:
            # Get confidence (max probability)
            d1_conf, _ = torch.max(d1_probs, dim=1)
            d2_conf, _ = torch.max(d2_probs, dim=1)
```

```

        # Average confidence
        confidence = (d1_conf + d2_conf) / 2

        if attention_weights is not None:
            return d1_pred, d2_pred, confidence, attention_weights
        else:
            return d1_pred, d2_pred, confidence

    else:
        if attention_weights is not None:
            return d1_pred, d2_pred, attention_weights
        else:
            return d1_pred, d2_pred

def get_num_params(self, trainable_only=True):
    """Return the number of parameters"""
    if trainable_only:
        return sum(p.numel() for p in self.parameters() if p.requires_grad)
    else:
        return sum(p.numel() for p in self.parameters())

if __name__ == "__main__":
    print("="*70)
    print("Testing Enhanced JerseyTemporalNet Model")
    print("="*70)

    # Test configuration
    batch_size = 4
    seq_length = 8
    img_size = 224
    num_classes = 11

    # Create model
    print("\n1. Creating model...")
    model = JerseyTemporalNet(
        num_classes=num_classes,
        hidden_dim=256,
        dropout=0.3,
        backbone='resnet18',

```

```

        use_spatial_attention=True,
        use_temporal_attention=True,
        bidirectional=True
    )

    print("    Model created!")
    print(f"    Total parameters:
{model.get_num_params(trainable_only=False):,} ")
    print(f"    Trainable parameters:
{model.get_num_params(trainable_only=True):,} ")

    # Create dummy input
    dummy_input = torch.randn(batch_size, seq_length, 3, img_size,
img_size)
    print("\n2. Testing forward pass...")
    print(f"    Input shape: {dummy_input.shape}")

    # Forward pass
    model.eval()
    output = model(dummy_input)

    if len(output) == 3:
        d1_logits, d2_logits, attention = output
        print(f"    Output D1 shape: {d1_logits.shape}")
        print(f"    Output D2 shape: {d2_logits.shape}")
        print(f"    Attention weights shape: {attention.shape}")
    else:
        d1_logits, d2_logits = output
        print(f"    Output D1 shape: {d1_logits.shape}")
        print(f"    Output D2 shape: {d2_logits.shape}")

    print("\n" + "="*70)
    print(" All tests passed!")
    print("="*70)

```

This is a neural network model for recognizing jersey numbers from video sequences.

**WHAT THIS MODEL DOES:**

It takes sequences of images (multiple frames showing a jersey number) and predicts two digits:

1. First digit (tens place) - can be 0-9 or "empty" (10) for single-digit numbers
2. Second digit (units place) - 0-9

MODEL ARCHITECTURE HAS 4 MAIN PARTS:

#### 1. SPATIAL FEATURE EXTRACTOR (CNN BACKBONE)

- Uses ResNet18/34/50 as the base
- Takes each image frame and extracts visual features
- Includes Spatial Attention module to focus on jersey number region
- Example: If frame shows a player, attention helps focus on the jersey number area

#### 2. TEMPORAL AGGREGATOR (LSTM)

- Takes features from all frames in the sequence
- Uses LSTM (Long Short-Term Memory) to understand temporal patterns
- Can be bidirectional - looks at frames both forward and backward in time
- Example: If number is partially visible in some frames, LSTM combines information across all frames

#### 3. TEMPORAL ATTENTION

- Weights different frames based on importance
- Some frames might be clearer than others
- Gives higher weight to frames where number is most visible
- Example: Frame where player is facing camera gets more weight than side view

#### 4. TWO CLASSIFICATION HEADS

- Head 1: Predicts first digit (tens place)
- Head 2: Predicts second digit (units place)
- Both heads use same features but have separate classifiers
- Example: For number "46": Head 1 predicts "4", Head 2 predicts "6"

KEY FEATURES:

Freezing Backbone Layers:

- ResNet starts with pretrained weights (trained on ImageNet)
- Freezes 60% of backbone layers initially (keeps them fixed)
- Allows fine-tuning only higher layers specific to jersey numbers
- Prevents overfitting when dataset is small

Spatial Attention:

- Learns to focus on relevant parts of the image
- Creates a "heatmap" highlighting jersey number region
- Helps ignore irrelevant background

Temporal Attention:

- Learns which frames are most informative
- Outputs weights showing importance of each frame
- Can see which frames the model considered important

Handling Variable Sequences:

- Uses masking for sequences with different lengths
- Ignores padding frames (empty frames added to make sequences same length)

TRAINING STRATEGY:

Two Separate Losses:

- Loss for first digit prediction
- Loss for second digit prediction
- Total loss = sum of both losses

Gradual Unfreezing:

- Starts with most backbone layers frozen
- Gradually unfreezes layers during training
- Allows model to adapt while keeping useful pretrained features

MODULE BREAKDOWN:

SpatialAttention class:

- Takes CNN features (e.g., 512 channels)
- Learns which spatial locations are important
- Multiplies features by attention weights

TemporalAttention class:

- Takes LSTM outputs for each frame
- Learns weights for each frame
- Combines frames using weighted average

EnhancedDigitHead class:

- Classification head for digit prediction
- Uses residual connections (helps training)
- Has dropout for regularization

JerseyTemporalNet class:

- Main model combining all components
- Has forward() method for training
- Has predict() method for inference

WORKFLOW EXAMPLE:

Input: 5 frames showing jersey number "23"

1. Each frame goes through ResNet → extract features
2. Spatial attention focuses on number region in each frame
3. LSTM processes all 5 frames together
4. Temporal attention weights frames (e.g., frame 3 gets 0.4 weight, frame 1 gets 0.1, etc.)
5. Weighted features go to two heads:
  - Head 1 predicts: "2" (tens digit)
  - Head 2 predicts: "3" (units digit)

#### OUTPUT FORMAT:

- Two sets of probabilities (11 classes each: 0-9 + "empty")
- Attention weights showing frame importance
- Confidence scores for predictions

#### TESTING SECTION:

The bottom code tests the model with dummy data to ensure:

- Model can be created
- Forward pass works
- Output shapes are correct
- Parameters are counted correctly

#### WHY THIS DESIGN:

- Uses temporal information: Multiple frames help when some frames are blurry/occluded
- Attention mechanisms: Focus on what matters (both spatially and temporally)
- Two-head design: Natural for two-digit numbers
- Pretrained backbone: Leverages general image knowledge
- Flexibility: Can handle different sequence lengths, different backbones

#### [train.py](#)

```
# train.py
import os
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, Subset
from torch.cuda.amp import autocast, GradScaler
from torch.utils.tensorboard import SummaryWriter
from torchvision import transforms
from tqdm import tqdm
import numpy as np
```

```

import json
import random
from datetime import datetime
import warnings
warnings.filterwarnings('ignore')

from config import Config
from dataset import JerseySequenceDataset
from model import JerseyTemporalNet


# ===== MIXUP =====
def mixup_data(x, d1_target, d2_target, alpha=1.0):
    if alpha > 0:
        lam = np.random.beta(alpha, alpha)
    else:
        lam = 1.0
    batch_size = x.size(0)
    index = torch.randperm(batch_size).to(x.device)
    mixed_x = lam * x + (1 - lam) * x[index]
    return mixed_x, d1_target, d1_target[index], d2_target,
d2_target[index], lam

def mixup_criterion(criterion, pred, y_a, y_b, lam):
    return lam * criterion(pred, y_a) + (1 - lam) * criterion(pred, y_b)
# =====


class EarlyStopping:
    def __init__(self, patience=10, min_delta=0.001, verbose=True):
        self.patience = patience
        self.min_delta = min_delta
        self.counter = 0
        self.best_loss = None
        self.early_stop = False
        self.verbose = verbose

    def __call__(self, val_loss):
        if self.best_loss is None:
            self.best_loss = val_loss

```



```

        self.best_val_acc = 0.0

    def setup_data(self):
        print("\n" + "="*60 + "\nSETTING UP DATA\n" + "="*60)

        self.train_transform = transforms.Compose([
            transforms.Resize((Config.IMG_SIZE, Config.IMG_SIZE)),
            transforms.RandomAffine(degrees=Config.AUG_ROTATION_DEGREES,
                                   translate=Config.AUG_TRANSLATE,
                                   scale=Config.AUG_SCALE),
            transforms.ColorJitter(brightness=Config.AUG_BRIGHTNESS,
                                   contrast=Config.AUG_CONTRAST),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225]),
        ])
        self.val_transform = transforms.Compose([
            transforms.Resize((Config.IMG_SIZE, Config.IMG_SIZE)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225])
        ])

        temp_ds = JerseySequenceDataset(Config.DATA_ROOT, transform=None,
                                         seq_length=Config.SEQ_LENGTH,
                                         allowed_classes=Config.TRAIN_CLASSES, mode='full')
        indices =
np.random.RandomState(Config.SEED).permutation(len(temp_ds))
        split = int(Config.VAL_SPLIT * len(indices))
        val_idx, train_idx = indices[:split], indices[split:]

        train_ds = Subset(JerseySequenceDataset(Config.DATA_ROOT,
self.train_transform,
                                         seq_length=Config.SEQ_LENGTH,
                                         allowed_classes=Config.TRAIN_CLASSES,
                                         mode='train',
                                         sampling_strategy=Config.TEMPORAL_SAMPLING),

```

```

        train_idx)
    val_ds = Subset(JerseySequenceDataset(Config.DATA_ROOT,
self.val_transform,
seq_length=Config.SEQ_LENGTH,

allowed_classes=Config.TRAIN_CLASSES,
mode='val',
sampling_strategy='uniform'),
val_idx)

full_train_ds = JerseySequenceDataset(Config.DATA_ROOT,
self.train_transform,
seq_length=Config.SEQ_LENGTH,

allowed_classes=Config.TRAIN_CLASSES, mode='train')
self.w_d1, self.w_d2 =
full_train_ds.get_class_weights(method='effective_number')

self.train_loader = DataLoader(train_ds,
batch_size=Config.BATCH_SIZE, shuffle=True,
num_workers=Config.NUM_WORKERS,
pin_memory=True,

persistent_workers=Config.NUM_WORKERS > 0)
self.val_loader = DataLoader(val_ds, batch_size=Config.BATCH_SIZE,
shuffle=False,
num_workers=Config.NUM_WORKERS,
pin_memory=True,
persistent_workers=Config.NUM_WORKERS
> 0)

print(f"Train batches: {len(self.train_loader)} | Val batches:
{len(self.val_loader)}")

def setup_model(self):
    print("\n" + "="*60 + "\nSETTING UP MODEL\n" + "="*60)
    self.model = JerseyTemporalNet(
        num_classes=Config.NUM_DIGIT_CLASSES,
        hidden_dim=Config.HIDDEN_DIM,
        dropout=Config.DROPOUT,

```

```

        backbone=Config.BACKBONE,
        use_spatial_attention=Config.USE_SPATIAL_ATTENTION,
        use_temporal_attention=Config.USE_TEMPORAL_ATTENTION,
        bidirectional=Config.BIDIRECTIONAL_LSTM
    ).to(self.device)

    self.criterion_d1 =
nn.CrossEntropyLoss(weight=self.w_d1.to(self.device), label_smoothing=0.1)
    self.criterion_d2 =
nn.CrossEntropyLoss(weight=self.w_d2.to(self.device), label_smoothing=0.1)

    self.optimizer = optim.AdamW(self.model.parameters(),
lr=Config.LEARNING_RATE, weight_decay=1e-4)
    self.scheduler =
optim.lr_scheduler.CosineAnnealingLR(self.optimizer,
T_max=Config.NUM_EPOCHS)

    # Log model graph to TensorBoard
    dummy_input = torch.randn(1, Config.SEQ_LENGTH, 3, Config.IMG_SIZE,
Config.IMG_SIZE).to(self.device)
    try:
        self.writer.add_graph(self.model, dummy_input)
        print(" Model graph saved to TensorBoard")
    except Exception as e:
        print(f" Could not save model graph: {e}")

    print("Model ready - MixUp + effective_number weights with
suppressed empty class")

def forward_with_output(self, x):
    output = self.model(x)
    if len(output) == 3:
        d1_logit, d2_logit, _ = output
    else:
        d1_logit, d2_logit = output
    return d1_logit, d2_logit

def calculate_metrics(self, d1p, d2p, d1t, d2t):
    both = ((d1p == d1t) & (d2p == d2t)).float().mean().item() * 100
    d1_acc = (d1p == d1t).float().mean().item() * 100

```

```

d2_acc = (d2p == d2t).float().mean().item() * 100
jersey_pred = torch.where(d1p == 10, d2p, d1p * 10 + d2p)
jersey_true = torch.where(d1t == 10, d2t, d1t * 10 + d2t)
jersey_acc = (jersey_pred == jersey_true).float().mean().item()
100
    return {'both_acc': both, 'd1_acc': d1_acc, 'd2_acc': d2_acc,
'jersey_acc': jersey_acc}

def train_epoch(self, epoch):
    self.model.train()
    total_loss = 0.0
    all_d1_pred, all_d2_pred, all_d1_true, all_d2_true = [], [], [], []

    for seq, d1_true, d2_true in tqdm(self.train_loader, desc=f"Epoch {epoch} [Train]"):
        seq = seq.to(self.device)
        d1_true = d1_true.to(self.device)
        d2_true = d2_true.to(self.device)

        self.optimizer.zero_grad()

        if random.random() < 0.5: # MixUp
            seq_mix, d1a, d1b, d2a, d2b, lam = mixup_data(seq, d1_true, d2_true, alpha=1.0)
            with autocast(enabled=self.use_amp):
                d1_logit, d2_logit = self.forward_with_output(seq_mix)
                loss = (mixup_criterion(self.criterion_d1, d1_logit, d1a, d1b, lam) +
                        mixup_criterion(self.criterion_d2, d2_logit, d2a, d2b, lam))
        else:
            with autocast(enabled=self.use_amp):
                d1_logit, d2_logit = self.forward_with_output(seq)
                loss = self.criterion_d1(d1_logit, d1_true) +
self.criterion_d2(d2_logit, d2_true)

        if self.use_amp:
            self.scaler.scale(loss).backward()
            self.scaler.unscale_(self.optimizer)

```

```

        torch.nn.utils.clip_grad_norm_(self.model.parameters(),
Config.GRADIENT_CLIP_VALUE)
            self.scaler.step(self.optimizer)
            self.scaler.update()
        else:
            loss.backward()
            torch.nn.utils.clip_grad_norm_(self.model.parameters(),
Config.GRADIENT_CLIP_VALUE)
            self.optimizer.step()

        total_loss += loss.item()
        all_d1_pred.append(d1_logit.argmax(1).cpu())
        all_d2_pred.append(d2_logit.argmax(1).cpu())
        all_d1_true.append(d1_true.cpu())
        all_d2_true.append(d2_true.cpu())

metrics = self.calculate_metrics(
    torch.cat(all_d1_pred), torch.cat(all_d2_pred),
    torch.cat(all_d1_true), torch.cat(all_d2_true)
)

avg_loss = total_loss / len(self.train_loader)

# Log to TensorBoard
self.writer.add_scalar('Loss/train', avg_loss, epoch)
self.writer.add_scalar('Accuracy/train_both', metrics['both_acc'],
epoch)
self.writer.add_scalar('Accuracy/train_d1', metrics['d1_acc'],
epoch)
self.writer.add_scalar('Accuracy/train_d2', metrics['d2_acc'],
epoch)
self.writer.add_scalar('Accuracy/train_jersey',
metrics['jersey_acc'], epoch)

return avg_loss, metrics

def validate(self, epoch):
    self.model.eval()
    total_loss = 0.0
    all_d1_pred, all_d2_pred, all_d1_true, all_d2_true = [], [], [], []

```

```

        with torch.no_grad():
            for seq, d1_true, d2_true in tqdm(self.val_loader, desc=f"Epoch {epoch} [Val] "):
                seq, d1_true, d2_true = seq.to(self.device),
                d1_true.to(self.device), d2_true.to(self.device)
                with autocast(enabled=self.use_amp):
                    d1_logit, d2_logit = self.forward_with_output(seq)
                    loss = self.criterion_d1(d1_logit, d1_true) +
                           self.criterion_d2(d2_logit, d2_true)

                    total_loss += loss.item()
                    all_d1_pred.append(d1_logit.argmax(1).cpu())
                    all_d2_pred.append(d2_logit.argmax(1).cpu())
                    all_d1_true.append(d1_true.cpu())
                    all_d2_true.append(d2_true.cpu())

    metrics = self.calculate_metrics(
        torch.cat(all_d1_pred), torch.cat(all_d2_pred),
        torch.cat(all_d1_true), torch.cat(all_d2_true)
    )

    avg_loss = total_loss / len(self.val_loader)

    # Log to TensorBoard
    self.writer.add_scalar('Loss/val', avg_loss, epoch)
    self.writer.add_scalar('Accuracy/val_both', metrics['both_acc'],
epoch)
    self.writer.add_scalar('Accuracy/val_d1', metrics['d1_acc'], epoch)
    self.writer.add_scalar('Accuracy/val_d2', metrics['d2_acc'], epoch)
    self.writer.add_scalar('Accuracy/val_jersey',
metrics['jersey_acc'], epoch)

    return avg_loss, metrics

def train(self):
    print("\n" + "="*60)
    print("TRAINING WITH MIXUP")
    print("=="*60)

```

```

        for epoch in range(1, Config.NUM_EPOCHS + 1):
            train_loss, train_metrics = self.train_epoch(epoch)
            val_loss, val_metrics = self.validate(epoch)
            self.scheduler.step()

            self.history['epoch'].append(epoch)
            self.history['train_loss'].append(train_loss)
            self.history['val_loss'].append(val_loss)
            self.history['train_acc'].append(train_metrics['both_acc'])
            self.history['val_acc'].append(val_metrics['both_acc'])
            self.history['val_d1_acc'].append(val_metrics['d1_acc'])
            self.history['val_d2_acc'].append(val_metrics['d2_acc'])

            self.history['val_jersey_acc'].append(val_metrics['jersey_acc'])

        self.history['learning_rate'].append(self.optimizer.param_groups[0]['lr'])

        # Log learning rate
        self.writer.add_scalar('Learning_Rate',
            self.optimizer.param_groups[0]['lr'], epoch)

        # Log train vs val comparison
        self.writer.add_scalars('Loss/train_vs_val', {
            'train': train_loss,
            'val': val_loss
        }, epoch)

        self.writer.add_scalars('Accuracy/train_vs_val', {
            'train': train_metrics['both_acc'],
            'val': val_metrics['both_acc']
        }, epoch)

        print(f"\nEpoch {epoch:2d} | Val Acc:
{val_metrics['both_acc']):5.2f}% | "
              f"D1: {val_metrics['d1_acc']):5.2f}% | D2:
{val_metrics['d2_acc']):5.2f}%")

        if val_metrics['both_acc'] > self.best_val_acc:
            self.best_val_acc = val_metrics['both_acc']
            torch.save({

```

```

        'epoch': epoch,
        'model_state_dict': self.model.state_dict(),
        'optimizer_state_dict': self.optimizer.state_dict(),
        'best_val_acc': self.best_val_acc,
        'history': self.history
    }, os.path.join(Config.CHECKPOINT_DIR, 'best_model.pth'))
    print(f"NEW BEST MODEL SAVED: {self.best_val_acc:.2f}"))

    if self.early_stopping(val_loss):
        print("Early stopping triggered")
        break

    # Close TensorBoard writer
    self.writer.close()

    print(f"\nTRAINING COMPLETE! Best Val Acc:
{self.best_val_acc:.2f}%")
    print(f" View TensorBoard: tensorboard --logdir={Config.LOGS_DIR}")
    print("Now run: python evaluate.py --checkpoint
checkpoints/best_model.pth")

if __name__ == "__main__":
    torch.manual_seed(Config.SEED)
    np.random.seed(Config.SEED)
    random.seed(Config.SEED)
    if torch.cuda.is_available():
        torch.cuda.manual_seed_all(Config.SEED)
        torch.backends.cudnn.deterministic = True
        torch.backends.cudnn.benchmark = False

    Trainer().train()

```

This is the main training script that trains the jersey number recognition model.

#### WHAT THIS SCRIPT DOES:

It trains the neural network using a combination of techniques to recognize jersey numbers from video sequences.

## MAIN COMPONENTS:

### 1. MixUp Augmentation:

- Creates new training samples by mixing two images
- Example: Mixes 30% of image with number "23" + 70% of image with number "45"
- Helps model generalize better and reduces overfitting
- Used randomly during training (50% chance per batch)

### 2. Early Stopping:

- Stops training if validation loss doesn't improve for 10 epochs
- Prevents overtraining
- Saves the best model

### 3. Trainer Class:

Manages the entire training process with these steps:

## SETUP PHASE:

### A. Data Preparation:

- Creates two different transforms:
  - Train transform: Adds augmentation (rotation, color jitter, etc.)
  - Val transform: Simple resize and normalize (no augmentation)
- Splits data into train (80%) and validation (20%)
- Calculates class weights to handle imbalanced data
- Creates data loaders for batching

### B. Model Setup:

- Creates the JerseyTemporalNet model
- Sets up loss functions with class weights
- Uses AdamW optimizer with weight decay
- Uses cosine annealing learning rate scheduler
- Sets up TensorBoard for visualization

## TRAINING LOOP:

For each epoch (training cycle):

### 1. Training Phase:

- Sets model to training mode
- For each batch:
  - Randomly applies MixUp (50% chance)
  - Passes sequences through model
  - Calculates loss for both digits
  - If MixUp was used: calculates combined loss from both mixed samples

- Backpropagates error
- Uses gradient clipping to prevent exploding gradients
- Updates model weights
- If mixed precision training is enabled: uses less memory for faster training
- Logs metrics to TensorBoard

## 2. Validation Phase:

- Sets model to evaluation mode
- No gradient calculation (faster)
- No augmentation
- Passes validation data through model
- Calculates accuracy metrics:
  - D1 accuracy (tens digit)
  - D2 accuracy (units digit)
  - Both digits correct
  - Jersey number accuracy (reconstructs full number from both digits)
- Logs metrics to TensorBoard

## 3. Model Saving:

- Saves model if validation accuracy improves
- Saves checkpoint with:
  - Model weights
  - Optimizer state
  - Training history
  - Best accuracy

## 4. Learning Rate Adjustment:

- Updates learning rate using cosine annealing
- Gradually decreases learning rate over time

## 5. Early Stopping Check:

- Checks if validation loss hasn't improved
- Stops training if no improvement for 10 epochs

## KEY METRICS CALCULATED:

### Accuracy Types:

1. D1 Accuracy: How often tens digit is correct (or correctly predicted as "empty" for single digits)
2. D2 Accuracy: How often units digit is correct
3. Both Digits Accuracy: Both D1 and D2 correct simultaneously
4. Jersey Number Accuracy: Full number reconstructed correctly
  - For single digits: matches units digit
  - For two digits: (D1 \* 10 + D2) matches actual number

## TRAINING TECHNIQUES USED:

### MixUp Augmentation:

- Randomly mixes two training samples
- Helps model learn smoother decision boundaries
- Reduces memorization of training data

### Class Weighting:

- Gives more weight to rare classes
- Prevents model from ignoring rare digits

### Label Smoothing:

- Makes labels slightly "fuzzy" (not exactly 1.0 for correct class)
- Helps prevent overconfidence
- Improves generalization

### Gradient Clipping:

- Limits gradient values
- Prevents exploding gradients during training

### Mixed Precision Training:

- Uses lower precision (FP16) where possible
- Faster training, less GPU memory
- Automatically managed by PyTorch

### Cosine Annealing Learning Rate:

- Learning rate follows cosine curve
- Starts high, decreases smoothly, may increase slightly at end
- Helps escape local minima

## TENSORBOARD VISUALIZATION:

- Logs training/validation losses
- Logs all accuracy metrics
- Logs learning rate changes
- Can visualize model graph
- View with: tensorboard --logdir=logs/

## OUTPUT:

- Best model saved as: checkpoints/best\_model.pth
- TensorBoard logs in: logs/run\_TIMESTAMP/
- Training history saved in checkpoint

## FINISHING:

- Prints best validation accuracy
- Suggests running evaluation script
- Closes TensorBoard writer

#### REPRODUCIBILITY:

- Sets random seeds for Python, NumPy, PyTorch
- Makes training deterministic (same results each run)

The script implements modern training techniques to create a robust model that can recognize jersey numbers even with limited training data, handling class imbalance and preventing overfitting.

#### [evaluate.py](#)

```
# evaluate.py

import os
import torch
import torch.nn as nn
from torch.utils.data import DataLoader
from torchvision import transforms
from tqdm import tqdm
import numpy as np
import json
from collections import defaultdict
from datetime import datetime
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

from config import Config
from dataset import JerseySequenceDataset
from model import JerseyTemporalNet


class Evaluator:
    def __init__(self, checkpoint_path):
        self.device = torch.device('cuda' if torch.cuda.is_available() else
'cpu')
        print(f"Using device: {self.device}")
        self.checkpoint_path = checkpoint_path
        self.load_checkpoint()
```

```

        self.setup_data()
        self.results = {'overall': {}, 'per_class': {}, 'predictions': [], 'confusion_matrices': {}}

    def load_checkpoint(self):
        print("\n" + "="*60)
        print("LOADING MODEL")
        print("=="*60)

        if not os.path.exists(self.checkpoint_path):
            raise FileNotFoundError(f"Checkpoint not found: {self.checkpoint_path}")

        checkpoint = torch.load(self.checkpoint_path,
map_location=self.device)

        if 'config' in checkpoint:
            cfg = checkpoint['config']
        else:
            # New format: vars(Config) saved directly
            cfg = {k: v for k, v in checkpoint.items() if k in [
                'num_digit_classes', 'hidden_dim', 'backbone',
'train_classes', 'test_classes'
            ]}

        # Reconstruct model
        self.model = JerseyTemporalNet(
            num_classes=cfg.get('num_digit_classes',
Config.NUM_DIGIT_CLASSES),
            hidden_dim=cfg.get('hidden_dim', Config.HIDDEN_DIM),
            dropout=Config.DROPOUT,
            backbone=cfg.get('backbone', Config.BACKBONE),
            use_spatial_attention=Config.USE_SPATIAL_ATTENTION,
            use_temporal_attention=Config.USE_TEMPORAL_ATTENTION,
            bidirectional=Config.BIDIRECTIONAL_LSTM
        ).to(self.device)

        self.model.load_state_dict(checkpoint['model_state_dict'])
        self.model.eval()

```

```

        print(f"Model loaded from: {self.checkpoint_path}")
        print(f"Train classes: {cfg.get('train_classes',
Config.TRAIN_CLASSES)}")
        print(f"Best val acc: {checkpoint.get('best_val_acc',
'N/A'):.2f}%)")
        print(f"Epoch: {checkpoint.get('epoch', 'N/A')} ")
        print(f"Parameters: {self.model.get_num_params():,} ")

    def setup_data(self):
        print("\n" + "="*60)
        print("SETTING UP TEST DATA")
        print("=".*60)

        self.test_transform = transforms.Compose([
            transforms.Resize((Config.IMG_SIZE, Config.IMG_SIZE)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225])
        ])

        self.test_dataset = JerseySequenceDataset(
            root_dir=Config.DATA_ROOT,
            transform=self.test_transform,
            seq_length=Config.SEQ_LENGTH,
            allowed_classes=Config.TEST_CLASSES,
            mode='test',
            sampling_strategy='uniform'
        )

        self.test_loader = DataLoader(
            self.test_dataset,
            batch_size=Config.BATCH_SIZE,
            shuffle=False,
            num_workers=Config.NUM_WORKERS,
            pin_memory=Config.PIN_MEMORY and self.device.type == 'cuda'
        )

        print(f"Test samples: {len(self.test_dataset)} | Batches:
{len(self.test_loader)}")

```

```

def evaluate(self):
    print("\n" + "="*60)
    print("RUNNING EVALUATION ON HELD-OUT CLASSES:",
Config.TEST_CLASSES)
    print("="*60)

    self.model.eval()
    all_d1_true, all_d2_true = [], []
    all_d1_pred, all_d2_pred = [], []
    all_d1_conf, all_d2_conf = [], []

    class_total = defaultdict(int)
    class_correct = defaultdict(int)

    # Anti-empty-bias heuristic parameters
    EMPTY_CLASS = 10
    EMPTY_THRESHOLD = 0.55 # If "empty" has < 55% confidence, use next
best

    total_modified = 0
    total_samples = 0

    with torch.no_grad():
        for seq, d1_true, d2_true in tqdm(self.test_loader,
desc="Evaluating"):
            seq = seq.to(self.device)
            d1_true = d1_true.to(self.device)
            d2_true = d2_true.to(self.device)

            # Safe forward (handles 2 or 3 return values)
            out = self.model(seq)
            if len(out) == 3:
                d1_logit, d2_logit, _ = out
            else:
                d1_logit, d2_logit = out

            # D2 prediction
            d2_pred = d2_logit.argmax(1)

```

```

# D1 prediction with anti-empty-bias heuristic
d1_probs_softmax = d1_logit.softmax(1)
d1_pred_list = []
batch_modified = 0

for i in range(len(d1_probs_softmax)):
    probs = d1_probs_softmax[i]

    if probs[EMPTY_CLASS] < EMPTY_THRESHOLD:

        masked_probs = probs.clone()
        masked_probs[EMPTY_CLASS] = 0
        prediction = masked_probs.argmax().item()
        d1_pred_list.append(prediction)

        # Count if we changed from what argmax would give
        if probs.argmax().item() == EMPTY_CLASS:
            batch_modified += 1
    else:

        d1_pred_list.append(EMPTY_CLASS)

d1_pred = torch.tensor(d1_pred_list,
device=d1_logit.device)

total_modified += batch_modified
total_samples += len(d1_pred_list)

# Get confidences for the predictions we made
d1_prob = d1_probs_softmax
d2_prob = d2_logit.softmax(1)
d1_conf = d1_prob.gather(1,
d1_pred.unsqueeze(1)).squeeze(1)
d2_conf = d2_prob.gather(1,
d2_pred.unsqueeze(1)).squeeze(1)

all_d1_true.extend(d1_true.cpu().numpy())
all_d2_true.extend(d2_true.cpu().numpy())
all_d1_pred.extend(d1_pred.cpu().numpy())
all_d2_pred.extend(d2_pred.cpu().numpy())

```

```

    all_d1_conf.extend(d1_conf.cpu().numpy())
    all_d2_conf.extend(d2_conf.cpu().numpy())

    # Per-class stats - CONFIDENCE-AWARE COMPOSITION
    for i in range(len(d1_true)):
        # True jersey number
        true_cls = int(d2_true[i].item() if d1_true[i] == 10
else d1_true[i].item()) 10 + d2_true[i].item()

        # Predicted jersey number (adaptive rule)
        if (d1_pred[i] != 10 and
            d1_conf[i] > Config.D1_CONF_THRESHOLD and
            d2_conf[i] > Config.D2_CONF_THRESHOLD):
            pred_cls = int(d1_pred[i].item() 10 +
d2_pred[i].item()) # two-digit
        else:
            pred_cls = int(d2_pred[i].item()) # fallback to
single-digit

        class_total[true_cls] += 1
        if pred_cls == true_cls:
            class_correct[true_cls] += 1

    print(f"\n Anti-empty-bias heuristic modified
{total_modified}/{total_samples} predictions " +
f"({100*total_modified/total_samples:.1f}%,\nthreshold={EMPTY_THRESHOLD})")

    # Convert to numpy
    d1_true = np.array(all_d1_true)
    d2_true = np.array(all_d2_true)
    d1_pred = np.array(all_d1_pred)
    d2_pred = np.array(all_d2_pred)
    d1_conf = np.array(all_d1_conf)
    d2_conf = np.array(all_d2_conf)

    # Overall accuracy (confidence-aware)
    overall_correct = 0
    for i in range(len(d1_true)):
```

```

        true_num = int(d2_true[i] if d1_true[i] == 10 else d1_true[i]
10 + d2_true[i])
        if (d1_pred[i] != 10 and
            d1_conf[i] > Config.D1_CONF_THRESHOLD and
            d2_conf[i] > Config.D2_CONF_THRESHOLD):
            pred_num = int(d1_pred[i] 10 + d2_pred[i]))
        else:
            pred_num = int(d2_pred[i])
        if pred_num == true_num:
            overall_correct += 1

    total = len(d1_true)
    acc = 100.0 * overall_correct / total
    d1_acc = 100.0 * (d1_true == d1_pred).mean()
    d2_acc = 100.0 * (d2_true == d2_pred).mean()
    avg_conf = np.mean((d1_conf + d2_conf) / 2)

    self.results['overall'] = {
        'accuracy': float(acc),
        'd1_accuracy': float(d1_acc),
        'd2_accuracy': float(d2_acc),
        'avg_confidence': float(avg_conf),
        'total_samples': int(total),
        'correct_samples': int(overall_correct),
        'empty_threshold': float.EMPTY_THRESHOLD,
        'predictions_modified': int(total_modified)
    }

    # Per-class
    for cls in sorted(class_total.keys()):
        self.results['per_class'][int(cls)] = {
            'total': class_total[cls],
            'correct': class_correct[cls],
            'accuracy': 100.0 * class_correct[cls] / class_total[cls]
        }

    # Confusion matrices (digit-wise only - reconstruction is post-hoc)
    self.results['confusion_matrices'] = {
        'd1': confusion_matrix(d1_true, d1_pred).tolist(),
        'd2': confusion_matrix(d2_true, d2_pred).tolist()
    }

```

```

        }

        self.print_results()
        self.save_results()
        self.generate_visualizations(d1_true, d2_true, d1_pred, d2_pred)

    def print_results(self):
        print("\n" + "="*60)
        print("GENERALIZATION TEST RESULTS (with Anti-Empty-Bias
Heuristic)")
        print("="*60)
        o = self.results['overall']
        print(f"Overall Accuracy: {o['accuracy']:.2f}%")
        print(f"Tens Digit Acc: {o['d1_accuracy']:.2f}%")
        print(f"Units Digit Acc: {o['d2_accuracy']:.2f}%")
        print(f"Avg Confidence: {o['avg_confidence']:.4f}")
        print(f"Total Samples: {o['total_samples']}")

        print("Predictions Modified:
{o['predictions_modified']}/{o['total_samples']}")

        print("\nPer-Class Accuracy:")
        for cls, stats in sorted(self.results['per_class'].items()):
            print(f"  Class {cls:2d} → {stats['accuracy']:6.2f}%
({stats['correct']}/{stats['total']})")

        if self.results['per_class']:
            best = max(self.results['per_class'].items(), key=lambda x:
x[1]['accuracy'])
            worst = min(self.results['per_class'].items(), key=lambda x:
x[1]['accuracy'])

            print(f"\nBest: Class {best[0]} → {best[1]['accuracy']:.2f}%")
            print(f"Worst: Class {worst[0]} → {worst[1]['accuracy']:.2f}%")

    def save_results(self):
        path = os.path.join(Config.RESULTS_DIR, 'evaluation_results.json')
        result_data = {
            'metadata': {
                'checkpoint': self.checkpoint_path,
                'test_classes': Config.TEST_CLASSES,
                'timestamp': datetime.now().isoformat(),

```

```

        'empty_threshold':
self.results['overall']['empty_threshold'],
        'reconstruction_rule': f"d1_conf >
{Config.D1_CONF_THRESHOLD} and d2_conf > {Config.D2_CONF_THRESHOLD}"
},
        'results': self.results
}
with open(path, 'w') as f:
    json.dump(result_data, f, indent=4)
print(f"\nResults saved to {path}")

def generate_visualizations(self, d1t, d2t, d1p, d2p):
    print("\nGenerating plots...")
    os.makedirs(os.path.join(Config.RESULTS_DIR, 'confusion_matrices'),
exist_ok=True)

    # Confusion matrices
    fig, ax = plt.subplots(1, 2, figsize=(16, 6))
    sns.heatmap(confusion_matrix(d1t, d1p), annot=True, fmt='d',
cmap='Blues', ax=ax[0])
    ax[0].set_title('Tens Digit Confusion')
    sns.heatmap(confusion_matrix(d2t, d2p), annot=True, fmt='d',
cmap='Greens', ax=ax[1])
    ax[1].set_title('Units Digit Confusion')
    plt.tight_layout()
    plt.savefig(os.path.join(Config.RESULTS_DIR, 'confusion_matrices',
'confusion_matrices.png'), dpi=300)
    plt.close()

    # Per-class bar chart
    if self.results['per_class']:
        classes = sorted(self.results['per_class'].keys())
        accs = [self.results['per_class'][c]['accuracy'] for c in
classes]
        plt.figure(figsize=(10, 6))
        bars = plt.bar(classes, accs, color=['green' if a >= 80 else
'orange' if a >= 60 else 'red' for a in accs])
        plt.axhline(y=self.results['overall']['accuracy'], color='red',
linestyle='--',

```

```

        label=f"Overall:
{self.results['overall']['accuracy']:.2f}%)"
    plt.title('Per-Class Accuracy on Test Set (with Anti-Empty-Bias
Heuristic)')
    plt.xlabel('Jersey Number')
    plt.ylabel('Accuracy (%)')
    plt.legend()
    plt.grid(axis='y', alpha=0.3)
    plt.savefig(os.path.join(Config.RESULTS_DIR,
'per_class_accuracy.png'), dpi=300, bbox_inches='tight')
    plt.close()

print("All plots saved!")

def main():
    import argparse
    parser = argparse.ArgumentParser()
    parser.add_argument('--checkpoint', type=str,
default='checkpoints/best_model.pth')
    args = parser.parse_args()

    evaluator = Evaluator(args.checkpoint)
    evaluator.evaluate()

    print("\n" + "="*60)
    print("EVALUATION COMPLETE!")
    print("=".*60)

if __name__ == "__main__":
    main()

```

This is the evaluation script that tests the trained model on new data to see how well it generalizes.

#### WHAT THIS SCRIPT DOES:

It loads a trained model and tests it on held-out jersey numbers (numbers the model hasn't seen during training) to evaluate generalization.

## MAIN STEPS:

### 1. Load the trained model:

- Loads checkpoint file saved during training
- Reconstructs the model architecture
- Loads the trained weights
- Prints model info (accuracy from training, parameters, etc.)

### 2. Prepare test data:

- Uses different jersey numbers than training (from Config.TEST\_CLASSES)
- Applies only basic transforms (no augmentation)
- Creates test data loader

### 3. Run evaluation with special heuristics:

The key innovation here is the "anti-empty-bias" heuristic and confidence-aware reconstruction.

#### Problem being solved:

The model tends to predict "empty" (10) for the tens digit too often, even for two-digit numbers. This script fixes that.

#### Anti-empty-bias heuristic:

- When model predicts "empty" for tens digit with low confidence (<55%)
- Ignore that prediction and choose the next most likely digit
- Example: Model predicts "empty" with 50% confidence, "4" with 45% confidence
- Heuristic chooses "4" instead of "empty"
- This fixes many wrong single-digit predictions

#### Confidence-aware reconstruction:

- Only reconstruct two-digit numbers if both digits have high confidence
- Uses thresholds: D1 confidence > 0.10, D2 confidence > 0.50
- Example: Model predicts "4" (confidence: 0.15) and "8" (confidence: 0.60)
- Both above thresholds → reconstructed number: 48
- If either digit has low confidence → use only units digit as single-digit number

### 4. Calculate multiple accuracy metrics:

- Overall accuracy: Full jersey number correct (using reconstruction rules)
- D1 accuracy: Tens digit correct (or correctly "empty")
- D2 accuracy: Units digit correct
- Per-class accuracy: Accuracy for each jersey number separately
- Average confidence: How confident the model is in its predictions

### 5. Generate visualizations:

- Confusion matrices: Show which digits get confused with which

- One for tens digit
- One for units digit
- Per-class bar chart: Shows accuracy for each jersey number
  - Green bars: Good accuracy ( $\geq 80\%$ )
  - Orange bars: Medium accuracy (60-80%)
  - Red bars: Poor accuracy (<60%)

## 6. Save results:

- Saves all metrics to JSON file
- Saves plots to results folder
- Includes metadata about the evaluation

## KEY FEATURES:

### Handling model outputs:

The script safely handles whether model returns 2 outputs (just logits) or 3 outputs (logits + attention weights)

### Per-class statistics:

Tracks accuracy for each jersey number separately to identify:

- Which numbers are easy/hard
- Patterns in errors

### Digit-level vs number-level accuracy:

- D1/D2 accuracy: Individual digit prediction
- Overall accuracy: Full number reconstructed correctly
- Jersey number accuracy: Using reconstruction rules

## PRACTICAL EXAMPLE:

Suppose test data has jersey number "48" (not in training, but digits "4" and "8" are in training numbers).

### Without heuristics:

- Model might predict: "empty" (10) for tens, "8" for units
- Reconstructed as: "8" (wrong)
- Reason: Model is biased toward predicting "empty" for tens digit

### With heuristics:

- Model predicts: "empty" with 50% confidence, "4" with 45% confidence, "8" with 85% confidence
- Anti-empty-bias: Choose "4" (not "empty") because empty confidence < 55%
- Confidence check: D1 confidence=0.45 > 0.10 ✓, D2 confidence=0.85 > 0.50 ✓
- Reconstructed as: "48" ✓

## OUTPUT FORMAT:

Console output shows:

- Overall accuracy percentage
- Individual digit accuracies
- Per-class accuracy breakdown
- Best/worst performing numbers
- How many predictions were modified by heuristics

Saved files:

- results/evaluation\_results.json (all metrics)
- results/confusion\_matrices/confusion\_matrices.png
- results/per\_class\_accuracy.png

## WHY THIS MATTERS:

This evaluation tests generalization:

- Train on numbers: 4, 6, 8, 9, 49, 66, 89
- Test on numbers: 48, 64, 88
- All test digits (4, 6, 8) appear in training, just in different combinations
- Tests if model learned to recognize individual digits, not just memorized specific numbers

## COMMAND LINE USE:

```

```
python evaluate.py --checkpoint checkpoints/best_model.pth
```

```

The script provides a comprehensive evaluation with diagnostic tools (confusion matrices, per-class breakdown) to understand exactly how the model performs and where it makes errors.

## analyze\_weights.py:

```
# analyze_weights.py
import torch
import json
import os
from config import Config
from model import JerseyTemporalNet

def analyze_model_weights(checkpoint_path):
    """Generate comprehensive weight summary"""

```

```
print("=="70)
print("TRAINED WEIGHTS SUMMARY")
print("=="70)

# Load checkpoint
checkpoint = torch.load(checkpoint_path, map_location='cpu')

print(f"\n Checkpoint: {checkpoint_path}")
print(f" Best Validation Accuracy: {checkpoint.get('best_val_acc',
'N/A'):.2f}%)")
print(f" Training Epoch: {checkpoint.get('epoch', 'N/A')}")

# Create model
model = JerseyTemporalNet(
    num_classes=Config.NUM_DIGIT_CLASSES,
    hidden_dim=Config.HIDDEN_DIM,
    dropout=Config.DROPOUT,
    backbone=Config.BACKBONE,
    use_spatial_attention=Config.USE_SPATIAL_ATTENTION,
    use_temporal_attention=Config.USE_TEMPORAL_ATTENTION,
    bidirectional=Config.BIDIRECTIONAL_LSTM
)

model.load_state_dict(checkpoint['model_state_dict'])

# Analyze weights by component
print("\n" + "=="70)
print("MODEL ARCHITECTURE SUMMARY")
print("=="70)

components = {
    'backbone': [],
    'spatial_attention': [],
    'lstm': [],
    'temporal_attention': [],
    'head_digit1': [],
    'head_digit2': [],
    'other': []
}
```

```

# Categorize parameters
for name, param in model.named_parameters():
    if 'backbone' in name:
        components['backbone'].append((name, param))
    elif 'spatial_attention' in name:
        components['spatial_attention'].append((name, param))
    elif 'lstm' in name:
        components['lstm'].append((name, param))
    elif 'temporal_attention' in name:
        components['temporal_attention'].append((name, param))
    elif 'head_digit1' in name:
        components['head_digit1'].append((name, param))
    elif 'head_digit2' in name:
        components['head_digit2'].append((name, param))
    else:
        components['other'].append((name, param))

# Print summary for each component
total_params = 0
trainable_params = 0

for component_name, params in components.items():
    if not params:
        continue

    component_total = sum(p.numel() for _, p in params)
    component_trainable = sum(p.numel() for _, p in params if
p.requires_grad)

    total_params += component_total
    trainable_params += component_trainable

    print(f"\n◆ {component_name.upper().replace('_', ' ')}")
    print(f"  Total Parameters: {component_total:,}")
    print(f"  Trainable: {component_trainable:,}")
    print(f"  ({100*component_trainable/component_total:.1f}%)")
    print(f"  Layers: {len(params)}")

# Show first 3 layers
for i, (name, param) in enumerate(params[:3]):

```

```

        print(f"      {name}: {list(param.shape)}")
if len(params) > 3:
    print(f"      ... ({len(params)-3} more layers)")

print("\n" + "="*70)
print("OVERALL STATISTICS")
print("="*70)
print(f"Total Parameters: {total_params:,}")
print(f"Trainable Parameters: {trainable_params:,}")
print(f"({100*trainable_params/total_params:.1f}%)")
print(f"Frozen Parameters: {total_params - trainable_params:,}")
print(f"Model Size: {total_params / (10242):.2f} MB (FP32)")

# Analyze digit heads specifically
print("\n" + "="*70)
print("DIGIT HEAD ANALYSIS (KEY FOR GENERALIZATION)")
print("="*70)

for head_name in ['head_digit1', 'head_digit2']:
    print(f"\n {head_name.upper()} (Tens Digit)" if '1' in head_name
else f"\n {head_name.upper()} (Units Digit)")
    head_params = components[head_name]

    for name, param in head_params:
        if 'fc_out.weight' in name:
            print(f"      Output Layer Shape: {list(param.shape)}")
            print(f"      → Maps {param.shape[1]} features →
{param.shape[0]} digit classes")
                print("      Weight Statistics:")
                print(f"          Mean: {param.data.mean().item():.4f}")
                print(f"          Std: {param.data.std().item():.4f}")
                print(f"          Min: {param.data.min().item():.4f}")
                print(f"          Max: {param.data.max().item():.4f}")

# Save summary to JSON
summary = {
    'checkpoint_path': checkpoint_path,
    'best_val_acc': float(checkpoint.get('best_val_acc', 0)),
    'epoch': int(checkpoint.get('epoch', 0)),
    'total_parameters': int(total_params),
}

```

```

'trainable_parameters': int(trainable_params),
'model_size_mb': float(total_params / (10242)),
'components': {
    name: {
        'total_params': int(sum(p.numel() for _, p in params)),
        'trainable_params': int(sum(p.numel() for _, p in params if p.requires_grad)),
        'num_layers': len(params)
    }
    for name, params in components.items() if params
}
}

output_path = os.path.join(Config.RESULTS_DIR, 'weight_summary.json')
with open(output_path, 'w') as f:
    json.dump(summary, f, indent=4)

print(f"\n Summary saved to: {output_path}")
print("=="*70)

if __name__ == "__main__":
    analyze_model_weights('checkpoints/best_model.pth')

```

This script analyzes the trained model's weights to understand how the model learned and what it learned.

#### WHAT THIS SCRIPT DOES:

It loads a trained model checkpoint and provides a detailed analysis of the model's internal weights - showing which parts of the model were most important, how many parameters were trained, and giving insights into what the model learned.

#### MAIN ANALYSIS STEPS:

##### 1. Load the checkpoint:

- Loads the saved model file
- Shows training accuracy and epoch number

##### 2. Create the model architecture:

- Recreates the exact same model structure

- Loads the trained weights into it

### 3. Categorize all parameters by component:

Groups weights into 7 categories:

- backbone: ResNet CNN layers (extract visual features)
- spatial\_attention: Layers that focus on jersey number region
- lstm: Temporal modeling layers
- temporal\_attention: Layers that weight different frames
- head\_digit1: Classifier for tens digit
- head\_digit2: Classifier for units digit
- other: Any remaining layers

### 4. Generate detailed statistics for each component:

For each component shows:

- Total parameters: How many numbers/weights
- Trainable parameters: How many were actually updated during training
- Percentage trainable: How much of this component was trained vs frozen
- Number of layers: How many distinct layers
- Example layer shapes: Shows the size of first few layers

### 5. Special focus on digit classification heads:

This is the most important analysis for understanding generalization:

Digit head output layer analysis:

- Shows weight matrix shape: [11 classes × number of features]
- For head\_digit1 (tens digit): Maps features → 11 classes (0-9 + "empty")
- For head\_digit2 (units digit): Maps features → 11 classes (0-9 + "empty")
- Shows weight statistics: mean, standard deviation, min, max

Why this matters:

The digit heads are what actually make the digit predictions. Their weights show:

- Which features are important for each digit
- How confident/uncertain the model is
- Whether it learned clear patterns or is confused

### 6. Calculate overall statistics:

- Total parameters: All weights in the model
- Trainable parameters: Weights that were updated
- Frozen parameters: Weights kept fixed (from pretrained ResNet)
- Model size: How much memory/disk space needed

### 7. Save results to JSON file:

- Saves all analysis to a file for future reference
- Includes component breakdown

- Includes training accuracy

#### KEY INSIGHTS YOU GET FROM THIS ANALYSIS:

##### 1. Understanding what was trained:

- Example: If backbone has 90% frozen, only 10% was fine-tuned
- This tells you if the model reused existing visual knowledge vs learned everything from scratch

##### 2. Parameter distribution:

- Which parts of model have most parameters
- Example: Backbone might have millions of parameters, heads only thousands
- Shows where model complexity is concentrated

##### 3. Training efficiency:

- Percentage of trainable parameters shows training efficiency
- Low percentage (e.g., 20% trainable) means leveraged pretrained knowledge well

##### 4. Digit head weight patterns:

- Weight magnitude: Large weights → strong feature importance
- Weight spread: High standard deviation → specialized features for different digits
- Extreme values: Very high/low weights → confident decisions

#### EXAMPLE OUTPUT INTERPRETATION:

If head\_digit2 output weights show:

- Mean: 0.001 (small weights overall)
- Std: 0.050 (moderate variation)
- Min: -0.100, Max: 0.150 (reasonable range)

This suggests:

- Model doesn't have extreme confidence
- Learned moderate distinctions between digits
- Should generalize reasonably

If instead weights showed:

- Mean: 0.500 (very large)
- Std: 2.000 (huge variation)
- Min: -5.000, Max: 8.000 (extreme values)

This would suggest:

- Overconfident model
- Possibly overfitted to training data
- Might not generalize well

## PRACTICAL USE CASES:

### 1. Model debugging:

- If accuracy is poor, check if backbone was properly unfrozen
- Check if digit heads have reasonable weight ranges

### 2. Model comparison:

- Compare different checkpoints to see training progress
- See which components changed most during training

### 3. Optimization:

- Identify components with most parameters for potential pruning
- See if certain heads could be made smaller without losing accuracy

### 4. Understanding generalization:

- Well-trained digit heads should show clear patterns
- Confused heads might show similar weights for all digits

## TECHNICAL DETAILS:

### Why analyze weights?

- Weights encode what the model learned
- Pattern in weights = pattern in learning
- Can reveal overfitting, underfitting, or good generalization

### Model size calculation:

- Each parameter is a 32-bit float (4 bytes)
- Total parameters  $\times$  4 bytes = size in bytes
- Divided by  $1024^2$  to get MB

### Frozen vs trainable:

- Frozen: From pretrained ResNet, kept unchanged
- Trainable: Updated during jersey number training
- Typically: Backbone mostly frozen, heads fully trainable

## COMMAND LINE USE:

Just run: `python analyze\_weights.py`

The script provides a "health check" for the trained model, showing what was learned, how complex the model is, and giving hints about whether it will generalize well to new data.